

Scheduling Real-Time Tasks in Multiprocessor and Distributed Systems

Real-Time Systems Design (CS 6414)

Sumanta Pyne

Assistant Professor
Computer Science and Engineering Department
National Institute of Technology Rourkela
pynes@nitrkl.ac.in

February 26, 2021

Overview

- Multiprocessor & distributed systems widely for real-time applications
- Recent price drop of these systems
- Dual processor architectures available at 50-60K INR even cheaper
- Distributed platform - networked PCs are common
- Better for faster response times and fault tolerance
- Distributed processing suitable for geographically distributed places
- Automated petroleum refinery - plant spread over geographic area
- Scheduling tasks more difficult than on a uniprocessor
- Uniprocessor - optimal schedule independent tasks in polynomial time
- Multiprocessor & Distributed - finding optimal schedule is NP-Hard

Multiprocessors & Distributed Systems

- Multiprocessors tightly coupled - shared physical memory
- Distributed Systems loosely coupled - no shared memory
- Interprocess communication (IPC) expensive in tightly-coupled syst.
- IPC time ignored compared to task execution times
- Inter-task communication - read/write on shared memory
- Distributed system IPC times comparable to task execution times
- Multiprocessor uses centralized scheduler/dispatcher
- Centralized scheduler maintain task state in centralized data structure
- High communicational overheads - update data structure state change

Scheduling on Multiprocessors & Distributed Systems

- Two subproblems
 - task allocation to processors
 - scheduling tasks on individual processors
- Task assignment problem
 - how to partition a set of tasks
 - then how to assign these to processors
 - Static - task allocation to nodes is permanent
 - Dynamic - tasks are assigned to nodes as they arise
 - different task instance allocated to different nodes
- Scheduling tasks on individual processors
 - Scheduling problem reduced to uniprocessor scheduling

Task allocation in Multiprocessors and Distributed Systems

- An NP-Hard Problem
- Exponential time required to determine an optimal schedule
- Usage of heuristic algorithms
- Task allocation algorithms - static and dynamic
- Static algorithms
 - all tasks partitioned into subsystems
 - each subsystem is assigned a separate processor
- Dynamic algorithms
 - tasks ready for execution placed in a common priority queue
 - dispatched to processors for execution as processors become available
 - different periodic task instances execute on different processors
- Most hard real-time systems till date are static in nature
- Dynamic real-time system make more efficient resource utilization

Multiprocessor Static Task Allocation

- Not applicable for distributed systems
- IPC time is same as memory access time due to shared memory
- Utilization Balancing algorithm
 - Maintains tasks in a queue in increasing order of utilizations
 - Removes tasks one by one from head
 - Allocates them to least utilized processor each time
 - To balance utilization of different processors
 - Perfectly balanced systems utilization u_i per processor equals overall
 - For task set ST_i assigned to processor P_i , $u_i = \sum_{t_j \in ST_i} u_{t_j}$
 - u_{t_j} utilization due to T_j
 - PR set of all processors, then total utilization is $\sum u_i$, $P_i \in PR$
 - Difficult to achieve perfect balancing of utilizations, i.e. $u_i = u'$ for P_i
 - Simple heuristic gives suboptimal results
 - Objective of good utilization balancing - Minimize $\sum |u' - u_i|$, $1 \leq i \leq n$
 - n #processors, u' average utilization of processors, u_i utilization of P_i
 - Suitable when #processors is fixed
 - Used when tasks at individual processors are scheduling using EDF

Next-Fit Algorithm for RMA

- A task set is partitioned, each scheduled on a uniprocessor using RMA
- Attempts to use as few processors as possible
- Classifies different tasks into few classes based on task utilization
- One or more processors assigned to each class of tasks
- Task with similar utilization values scheduled on same processor
- Policy for utilization based task classification
 - If tasks divided to m classes, $T_i \in \text{class } j, 0 \leq j \leq m$
- Partitioning tasks of a system into four classes
 - Class 1: $(2^{\frac{1}{2}} - 1) < C_1 \leq (2^{\frac{1}{1}} - 1)$
 - Class 2: $(2^{\frac{1}{3}} - 1) < C_2 \leq (2^{\frac{1}{2}} - 1)$
 - Class 3: $(2^{\frac{1}{4}} - 1) < C_3 \leq (2^{\frac{1}{3}} - 1)$
 - Class 4: $0 < C_4 \leq (2^{\frac{1}{4}} - 1)$
- Utilization grid for different classes
 - class 1:(0.41,1), class 2:(0.26,0.41), class 3:(0.19,0.26), class 4:(0,0.19)
- Higher task utilization values are coarser compared to lower
- Grid size: class 1 tasks is $1 - 0.41 = 0.59$, class 3 tasks is 0.7
- Simulation - next-fit algorithm atmost $2.34 \times$ optimum #processors

Example 1

The following table shows the execution times (in msec) and periods in (msec) of a set of 10 periodic real-time tasks. Assume that the tasks need to run on a multiprocessor with four processors. Allocate the tasks to processors using the next fit algorithm. Assume that the individual processors are to be scheduled using RMA algorithm.

Task Set

Task	T_1	T_2	T_3	T_4	T_5	T_6	T_7	T_8	T_9	T_{10}
e_i	5	7	3	1	10	16	1	3	9	17
p_i	10	21	22	24	30	40	50	55	70	100

Solution

Task	T_1	T_2	T_3	T_4	T_5	T_6	T_7	T_8	T_9	T_{10}
e_i	5	7	3	1	10	16	1	3	9	17
p_i	10	21	22	24	30	40	50	55	70	100
u_i	0.5	0.33	0.14	0.04	0.33	0.4	0.02	0.05	0.13	0.17
Class	1	2	4	4	2	2	4	4	4	3

Bin Packing for EDF

- Allocates tasks to processors
- Tasks on individual processors scheduled using EDF
- Tasks assigned to P_i such that $u_i \leq 1$
- Formulate task allocation as bin packing problem
 - Given n periodic real-time tasks
 - Individual processors scheduled using EDF
 - #bins necessary = $\sum u_i, 1 \leq i \leq n$
 - Bin packing is NP-complete
- First-fit random bin packing algorithm
 - Tasks selected randomly
 - Assigned processors arbitrarily as long as $u_i \leq 1$
 - Atmost $1.7 \times$ optimum #processors required
- First-fit decreasing bin packing algorithm
 - Tasks sorted in non-increasing order of CPU utilization in ordered list
 - Task selected one by one from ordered list
 - Assigned to bin(processor P_i) to which it fits first (i.e. $u_i \leq 1$)
 - Simulations show #processors required is $1.22 \times$ optimal #processors

Dynamic Allocation of Tasks

- Applications where tasks arrive sporadically at different nodes
- Dynamic algorithms needed to handle such tasks
- Assume any task can be executed on any processor
- Dynamic solutions naturally distributed
- No central allocation policy running on some processor
- No preallocation of tasks to processors
- Assign tasks when they arise
- Task allocation made on instantaneous load positions of nodes
- Achievable schedulable utilization better than static approaches
- High run time overhead - allocator component running at each node
- Static allocation - task assigned permanently during initialization
- No runtime overhead for static
- Dynamic ineffective if task bound to single or subset of processors

Focussed Addressing and Bidding

- Every processor maintain status table and system load table
- Status table contains
 - information of tasks committed to run
 - execute time of tasks
 - periods of tasks
- System load table contains latest load information other processors
- Determine surplus computing capacity available
- Time axis divided into windows - intervals of fixed duration
- At end of each window each processor broadcasts to all other
- Fraction of computing power currently free for next window
- Fraction of next window with no committed tasks
- On receiving this every processor updates system load table
- When task arises, node checks whether process it locally
- If possible, it updates its status table
- Otherwise, looks out for a process can offload task

Focussed Addressing and Bidding

- Processor finds suitable processor consulting system load table
- Determines least loaded processors to accommodate the task
- It then sends request for bids (RFBs) to these processors
- While looking for processor an overloaded processor
- Checks surplus information and selects a *focussed* processor
- information in system load table may be out of date
- Obsolete information problem
- Solution - send RFBs only if it determines that task complete in time
- High communication overhead to maintain system load tables
- Larger window size lower overhead

Buddy Algorithm

- Aims to overcome high communication overhead
- Processor states - underloaded and overloaded
- P_i underloaded if $u_i < Th$
- P_i overloaded if $u_i \geq Th$
- Processor broadcasts - change in state
- Broadcasts to limited subset *buddy* set

Fault-Tolerant Scheduling of Tasks

- Task scheduling can be used for fault-tolerance in real-time systems
- An effective technique
- Requires very little redundant hardware resource
- Achieved by scheduling additional ghost copies with primary task copy
- Ghost copies (clones)
 - may not be identical to primary copy
 - copies stripped down versions
 - can be executed in shorter durations than primary
- Ghost copies of different tasks can be overloaded on same slot
- Success execution of primary, deallocates corresponding backup

Clocks in Distributed Real-Time Systems

- Clocks used for determining timeouts and time stamping
- Timeouts
 - determine failure of tasks due to deadline miss
 - indicate transmission faults or delays, or non-existent receivers
- Time stamping
 - used in message communication among tasks
 - sender includes current time along a message
 - gives age of message
 - used for ordering purposes
 - relies on good real-time clock services
- Distributed system has one clock at each node
- Different clocks diverge
- Two clocks of run exactly at same speed is impossible

Clocks in Distributed Real-Time Systems


- Lack of synchrony among clocks expressed as *clock slew*
- *clock slew* determines attendant drift of clocks with time
- Lack of synchrony and drift - meaningless time stamping and timeout
- Synchronized clocks for meaningful time stamping and timeout
- Synchronization very important in distributed real-time systems

Clock Synchronization

- Makes all clocks in network agree on their time values
- Different clocks agree on some time, differ from world time standard
- World time standard called Universal coordinated time (UTC)
- UTC based on international atomic time (TAI)
- TAI maintained at Paris by averaging atomic clocks around the world
- UTC signals used through GPS receivers and specialized radio stations
- Internal clock synchronization - all clocks synchronized wrt one clock
- External clock sync - a set of clocks synchronized with external clock
- Two main approaches for internal clock synchronization
 - centralized clock synchronization
 - distributed clock synchronization

Centralized Clock Synchronization

- One of the clocks is designated as master clock or *time server*
- Other clocks of the system are slaves
- Slaves are kept in synchronization with master clock
- Slave clocks C_1, \dots, C_n are to synchronized with master clock



- Server broadcasts its time to all clocks after every ΔT time interval
- On receiving a broadcast, slaves set clock as per time at master clock
- ΔT should be chosen carefully
- If ΔT too small
 - frequent broadcasts from master
 - good synchronization between slaves and master
 - high communication overhead
- If ΔT too large, clocks may drift too much apart

Centralized Clock Synchronization

- Assume maximum rate of drift between to individual clocks is ρ
- Clock manufactures provides ρ as a specification parameter of a clock
- ρ is unit less, it measures drift (time) per unit time
- Suppose clocks are resynchronized after every ΔT interval
- Ignoring communication time required for broadcast
- Once broadcast received, clocks set to received time instantly
- Drift of any clock from master clock is bounded by $\rho\Delta T$
- Maximum drift between any two clocks is limited to $2\rho\Delta T$
- In reality it takes a finite amount of time to set a clock
- Suitable communication time and clock setting time required
- Otherwise, synchronized time become slower wrt external clock
- Slaves still remain synchronized within a specified bound
- Very difficult to compensate these in practical systems

Example 2

Assume that the drift rate between any two clocks is restricted to $\rho = 5 \times 10^{-6}$. Suppose we want to implement a synchronized set of six distributed clocks using the central synchronization scheme so that the maximum drift between any two clocks is restricted to $\epsilon = 1$ msec at any time, determine the period with which the clocks need to be resynchronized.

- Solution. The maximum drift rate between any two arbitrary clocks when the clocks are synchronized using a central time server with a resynchronization interval of ΔT is given by $2\rho\Delta T < \epsilon$. Therefore, the required resynchronization interval ΔT can be expressed as:


$$\Delta T < \frac{1 \times 10^{-3}}{5 \times 10^{-6} \times 2} \text{ sec} = \frac{10^{-3}}{10^{-5}} \text{ sec} = \frac{1}{10^{-2}} \text{ sec} = 100 \text{ sec.}$$

Therefore, resynchronization period must be less than 100 sec.

Distributed Clock Synchronization


- Centralized clock synchronization susceptible to single point failure
- Any failure of master clock causes breakdown of synchronization
- Distributed clock synchronization - no master clock wrt all slaves
- All clocks periodically exchange clock readings among themselves
- Based on received time, clocks compute and set synchronized time
- Possible that some clocks are bad or become bad during operation
- Bad clocks exhibit larger drifts than specified tolerance
- Bad clocks may even stop keeping time

Distributed Clock Synchronization



- Bad clocks can be identified and taken care of during synchronization
- By rejecting time values of any clock larger than specified bound
- Usage of Byzantine clocks
- A Byzantine clock is a two face clock
- Transmits different values to different clocks at same time

Byzantine Clock



- C_1 is a Byzantine clock
 - sending time value $t+e$ to clock C_5
 - $t-e$ to clock C_3
 - at same time instant
- If $< \frac{1}{3}$ clocks bad (Byzantine), good clocks approx. synchronized
- Synchronization scheme of clocks
 - Let there be n clocks in a system
 - Each clock periodically broadcasts time value at end of certain interval
 - Assume clocks to be synchronized within ϵ time units of each other
 - If a clock receives broadcast time differs from its own time by $> \epsilon$
 - then sending clock must be bad, safely ignores received time value
 - Each clock averages all good time values received
 - sets its time with average

Pseudo Code for Distributed Clock Synchronization

- Each clock C_i carries out following operations


Procedure distributed clock synchronization:

```
good-clocks=n;  
for(j=1;j < n;j++){  
    if ( $|(c_i - c_j)| > \epsilon$ ) good-clocks--; // Bad clock  
    else total-time= total-time +  $c_j$ ;  
     $c_i$ =total-time/good-clocks; // set own time equal to the computed time  
}
```

- Each clock carries same set of steps
- If all n clocks carry these steps, then atmost m are bad
- $n > 3m$
- Good clocks will be synchronized within $3\epsilon m/n$ bound

Theorem 1

In a distributed system with n clocks, a single Byzantine clock can make two arbitrary clocks in a system to differ by $3\epsilon/n$ in time value, where ϵ represents the maximum permissible drift between two clocks.



- Proof. Consider three clocks C1, C2 and C3 in a distributed system. C1 is a Byzantine clock. C2 and C3 are good clocks - differ by $\neq \epsilon$. C3 being Byzantine shows two different values to C1 and C2. Effect of Byzantine clock in total time calculation makes two good clocks differ by at most 3ϵ . Effect of a single Byzantine clock make two arbitrary clocks differ by $3\epsilon/n$. For m Byzantine clocks two good clocks differ by at most $3\epsilon m$ in average. Individual clocks synchronized within $3\epsilon m/n$.

Time required for clocks drift to ϵ

- Let time required for two clocks to drift from $3\epsilon m/n$ to ϵ be ΔT .
or, $2\Delta T\rho \leq (n\epsilon - 3\epsilon m)/n$
or, $\Delta T \leq (n\epsilon - 3\epsilon m)/(n \times 2\rho)$
 ΔT is time required for two good clocks to drift from $3\epsilon m/n$ to ϵ
 $\Delta T \leq [(3m+1)\epsilon - 3\epsilon m]/(n \times 2\rho)$
 $\Delta T \leq \epsilon/2n\rho$

Example 3

Let a distributed real-system have 10 clocks, and it is required to restrict their maximum drift to $\epsilon = 1$ msec. Let the maximum drift of the clocks per unit time (ρ) be 5×10^{-6} . Determine the required synchronization interval.

- Solution. $\Delta T = \frac{10^{-3}}{2 \times 10 \times 5 \times 10^{-6}}$

$$\Delta T = 10 \text{ sec}$$

Required synchronization interval is 10 sec.

Thank you